

Background

The TD partitioning VMM was implemented fully in the COCONUT-SVSM kernel mode, which runs the COCONUT-SVSM as a L1 VMM in TD and a Linux OS as a L2 guest. The code is published here. Following the COCONUT-SVSM development plan, most parts of the TD partitioning VMM should be moved to the user mode for security reasons, while only keeps the necessary parts in the kernel.

Similar for the SEV, the protocol requests from SEV-SNP guest now are handled by COCONUT-SVSM kernel, which should also be moved to the user mode.

So this document will describe a straightforward implementation for the user mode VMM utilizing the syscalls outlined in the X86-64 Syscall API/ABI Document.

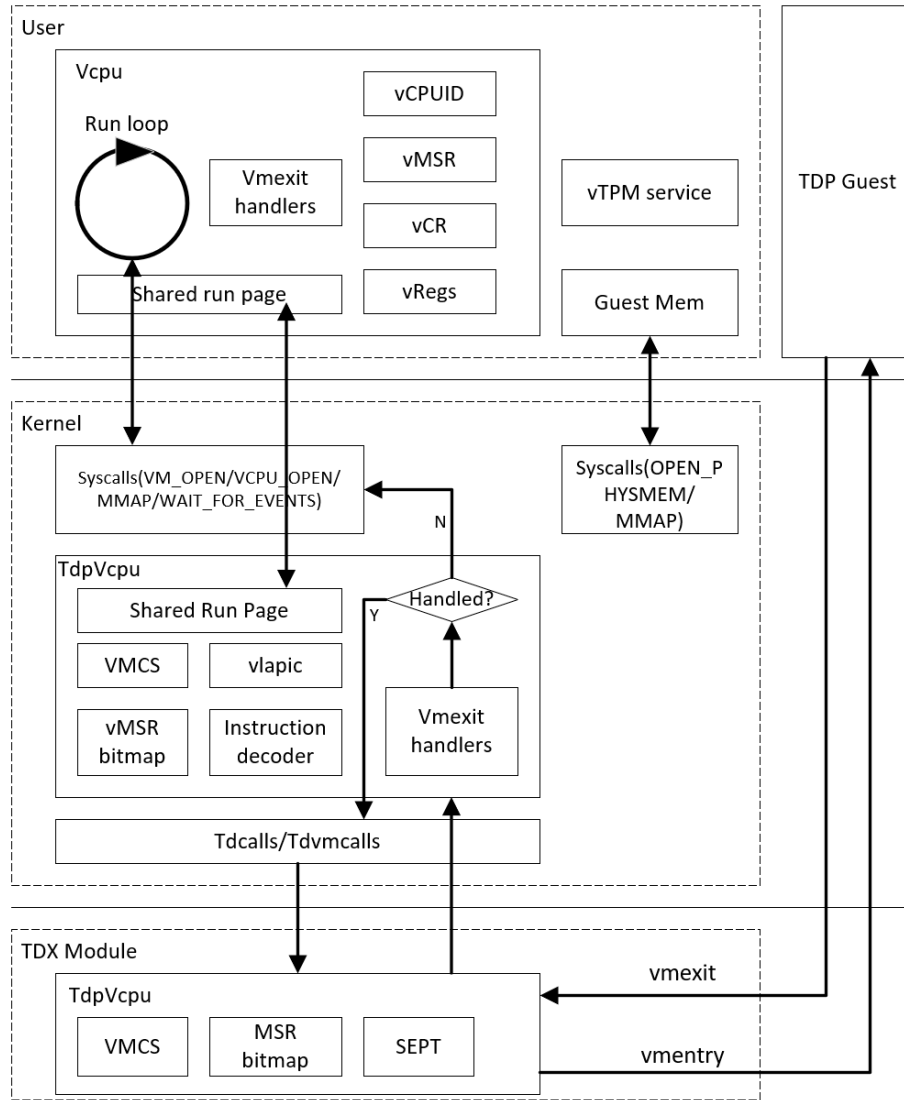
TDX User Mode Restrictions

According to Intel TDX CPU Architecture Extension Section 2.3 Instruction Set Reference TDCALL, the user mode is not allowed to execute the tdcall instruction as this can cause #GP when CPL > 0. Due to this restriction, any operation done via tdcall cannot be performed in the user mode, e.g.:

- Accessing VMCS, modifying MSR bitmap, configuring L2 VM control (TDG.VP.RD/WR)
- Entering L2(TDG.VP.ENTER)
- Adding page alias(TDG.MEM.PAGE.ATTR.WR)
- All kinds of tdvmcall with host VMM(TDG.VP.VMCALL)

VMM High Level Architecture

With the above user mode restriction for TDX, the tdcall instruction related parts are kept in the COCONUT-SVSM kernel. While considering the performance, vLAPIC and instruction decoder are also kept in the COCONUT-SVSM kernel. Below is a high level architecture for the VMM:



In this architecture, the user mode VMM responds to create vcpu threads, which each thread maintains a vcpu instance that contains the vcpu contexts (e.g., cpuid/msr/control registers/general purpose registers). The vcpu thread runs the loop of entering guest and handling guest vmexit events. The user mode VMM also responds to provide the secure device emulation, e.g., vTPM.

The COCONUT-SVSM kernel responds for tdcall related functionalities, e.g.:

- Access VMCS
- Modify the MSR bitmap.
- Inject interrupt.

- Send `tdvmcalls` to host VMM(e.g. IO/MMIO requests, page conversion).
- Do `vmentry`.

The COCONUT-SVSM kernel records the detailed `vmexit` information in a page shared with the user mode VMM. This allows the `vcpu` thread to retrieve the `vmexit` information from the shared page once it returns to the user mode for handling the `vmexit` events. But as the COCONUT-SVSM kernel has certain functionalities aforementioned, some `vmexit` reasons can be handled directly in the COCONUT-SVSM kernel without switching back to the user mode. For example:

- The EPT violation caused by missing page alias mapping in secure EPT.
- The virtual LAPIC MMIO/MSR `vmexit`.

The instruction decoding will be done in the kernel mode for performance reasons. Taken `vTPM` MMIO read instruction as an example, the kernel thread decodes the MMIO read instruction to get detailed MMIO address/data, and then wakes up the the user mode `vTPM` thread which is waiting on `vTPM` MMIO event to emulate. Once this is done, the user mode `vTPM` thread wakes up the kernel thread and waits for `vTPM` MMIO event again. Then the kernel thread completes the instruction emulation (writing back the emulated result) and enter the guest.

User Mode VMM

The user mode VMM is used to launch guest VM, and handle various `vmexit` events with necessary support from the COCONUT-SVSM kernel. It is compiled as a dedicated ELF, wrapped in the SVSM ram file system binary, and launched by the `init` process.

Main Thread

The user mode VMM starts to run at the main thread. It queries the valid VM bitmap via the `CAPABILITIES` syscall, and gets the `VmObjHandle` for the first valid VM from the COCONUT-SVSM kernel via `VM_OPEN` syscall. With the `VmObjHandle`, the user mode VMM can get the supported number of vCPUs via the `VM_CAPABILITIES` syscall, so that the main thread can create dedicated `vcpu` thread for each `vcpu` via the `THREAD_CREATE` syscall.

```
struct ObjHandle(i32); // A raw object handle
struct VmObjHandle(ObjHandle); // Represents vm object handle
```

Create Vm Structure

The main thread creates a `Vm` structure to represents an opened VM, which contains a `VmObjHandle` get via the `VM_OPEN` and the number of `vcpu` via the `VM_CAPABILITIES`. Dropping the structure `Vm` should result in calling the `CLOSE` syscall to close its `VmObjHandle`.

```

struct Vm {
    handle: VmObjHandle, // vm object handle
    nr_vcpus: u64, // supported number of vCPUs
    <!-- More fields can be extended in the future -->
}

```

Create Vcpu Thread

The main thread will get a `VcpuObjHandle` via the `VCPU_OPEN` syscall and start a dedicated vcpu thread via the `THREAD_CREATE` syscall for each vcpu.

The `THREAD_CREATE` takes `X86Regs` as the input parameter, which is used as the initial context for the new created thread. The `rip` should be set with the entry point. To allow passing one input parameter to the entry point, the `rdi` should also be set. The created thread should share certain user mappings with the parent(caller), like text/data sections, heap area, but not the stack area. This will require the `struct task` in kernel to have two separate user VMR, one is shared virtual memory range, which maps the text/data/heap and the other is private virtual memory range, which maps the stack. The parent task may need to record which tasks has been created, and a child task may also need to know which task is the parent. This is useful when terminate tasks. After a thread being created, the parent can use `THREAD_JON` syscall to wait for the child task finishing execution. So the `struct task` may require a `WaitQueue` for the parent task to wait. Below shows the new fields for `struct task` in the kernel mode:

```

struct task {
    <!-- Original fields -->
    /// Task shared virtual memory range for use at CPL 3 - None for kernel tasks
    vm_user_shared_range: Option<Arc<VMR>>,
    /// Task private virtual memory range for use at CPL 3 - None for kernel tasks
    vm_user_private_range: Option<VMR>,
    /// Child threads created by this task
    threads: RWLock<BTreeMap<u32, Arc<Task>>>,
    /// Parent task ID
    parent_id: Option<u32>,
    /// Wait queue to hold pending tasks which are joining this task
    join_wq: SpinLock<WaitQueue>,
}

```

For the vcpu thread, it needs to take the address of the `VcpuObjHandle` as the input parameter for its entry point so that the vcpu thread can retrieve the `VcpuObjHandle` opened by the main thread. As it shares the user mappings text/data/heap with the main thread except for the stack, the `VcpuObjHandle` should be put in the heap by the main thread to make sure its address can be accessed by the vcpu thread.

```

struct VcpuObjHandle(ObjHandle); // Represents vcpu object handle

```

After all the vcpu threads being created, the main thread can use the `THREAD_JOIN` syscall to wait for each vcpu thread to finish its execution, if it doesn't need to wait for any other event.

Vcpu Thread

The vcpu thread responses for running a vcpu via `WAIT_FOR_EVENT` syscall, and handling various vmexit events after being waked up. The detailed vmexit event information is stored in a 4KiB page provided by the vcpu object, which can be mapped to the vcpu thread's address space via `MMAP` syscall.

Create Vcpu Structure

A Vcpu structure is created by the vcpu thread to represent an opened VCPU. It contains the `VcpuObjHandle` which is the input parameter of the entry point, and a vcpu run page which points to the mapped 4KiB page provided by the vcpu object. Dropping the Vcpu structure should result in `munmap` the vcpu run page via `MUNMAP` syscall and close the opened vcpu object via `CLOSE` syscall.

```
struct VcpuRunPage(i64); // Represents vcpu run page address via MMAP syscall.

struct Vcpu {
    handle: VcpuObjHandle, // vcpu object handle
    run_page: VcpuRunPage, // vcpu run page
    <!-- More fields can be extended in the future -->
}
```

Vcpu Run Page

The vcpu object in the kernel provides a 4KiB page, which is used as a shared memory between the user mode and the kernel mode for communicating the vmexit events at the runtime. Typically, the kernel mode can write detailed vmexit event from the guest to this page before switching to the user mode vcpu thread. The user mode vcpu thread can get the detailed vmexit event from this page and update the handling result to this page before entering the kernel mode to do vmenter.

The `MMAP` syscall can map this 4KiB vcpu run page into the user private virtual memory range of the vcpu thread, as this is a per-vcpu page and no need to be shared with the main thread or any other vcpu thread. A `i64` is returned to the user mode as the virtual address, which can be used to read from or write to this 4KiB page.

The structure `VcpuRunPage` is defined to represent this 4KiB page, which contains an `Option` to represent a vmexit event.

```
enum VcpuExitEvent {
    // Cpuid vmexit event
```

```

Cpuid {
    leaf: u32,
    subleaf: u32,
    eax: u32,
    ebx: u32,
    ecx: u32,
    edx: u32,
},
// Protocol request event
ProtocolReq {
    protocol: u32,
    request: u32,
    sev_features: u64,
    rcx: u64,
    rdx: u64,
    r8: u64,
},
<!-- More fields can be extended in the future -->
}

#[repr(align(4096))]
struct VcpuRunPage {
    exit_event: Option<VcpuExitEvent>, // Guest vmexit event
    <!-- More fields can be extended in the future -->
}

```

Run Vcpu

The vcpu thread in the user mode cannot directly do vmenter. Instead, vm entering is done by a per-cpu kernel thread which is created during COCONUT-SVSM kernel booting time. After this per-cpu kernel thread being created, it starts to wait until being waked up to do vmenter.

So to request vmenter by a user mode vcpu thread, it can use the `WAIT_FOR_EVENT` syscall, which takes the `ObjHandle` (raw of the `VcpuObjHandle`) as the input parameter. For the vcpu object, syscall will make the vcpu thread waiting on the event belongs to the vcpu object, and wakes up the corresponding per-cpu kernel thread to do vmenter. After vmexit happens from the guest, the per-cpu kernel thread wakes up the vcpu thread and waits again. When vcpu thread backs to the user mode, it can retrieve the vmexit event from the vcpu run page and do the handling.

Basic Flow Overview

With the aforementioned syscalls, the user mode VMM is able to run vcpu and handle certain vmexit events. Below is an overview of this basic flow:

